MICAI 25 YEARS 2000-2026
02-06 NOV 2026
MEXICAN INTERNATIONAL CONFERENCE ON ARTIFICIAL INTELLIGENCE
|ITESM Campus Chihuahua

Paper #12 · Oral

Forecasting Weekly Depression Incidence in Mexico:
A Multi-Model Framework with Auditable Per-Series Model Selection

Javier Augusto Rebull-Saucedo[1] · **Juan Carlos Pérez-Nava**[1] · **Luis Gerardo Sánchez-Salazar**[1]
Grettel Barceló-Alonso[1] · **Ruth Pérez-Hernández**[2]

[1] Tecnológico de Monterrey, Monterrey, Mexico
[2] Instituto Mexicano del Seguro Social, OOAD Guerrero, Mexico

When recorded 2026 forecasts meet observed cases

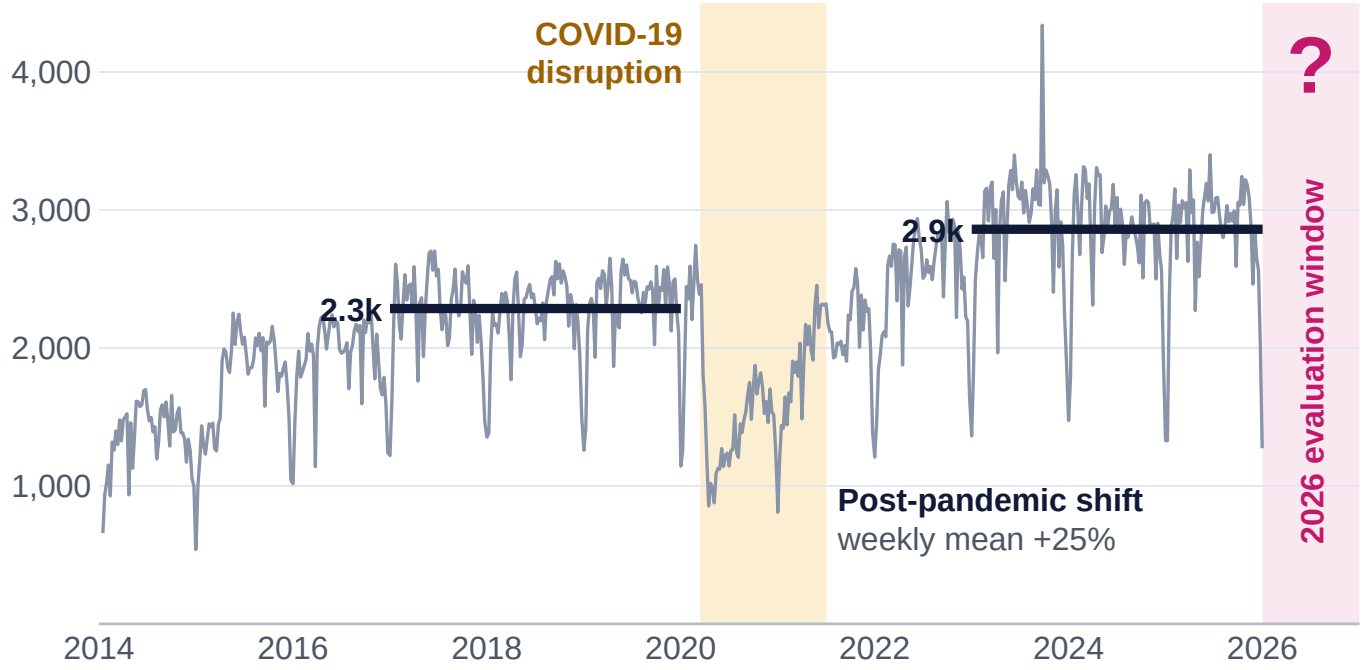


Sociedad Mexicana de Inteligencia Artificial
Tecnológico de Monterrey | CAMPUS CHIHUAHUA 50 AÑOS
AMBASSADE DE FRANCE AU MEXIQUE
MÉXICO FRANCIA 2026 200 AÑOS DE AMISTAD
CHIHUAHUA BURÓ DE CONVENCIONES
¡ah Chihuahua!

01 THE PROBLEM

Depression is counted weekly, but only in hindsight

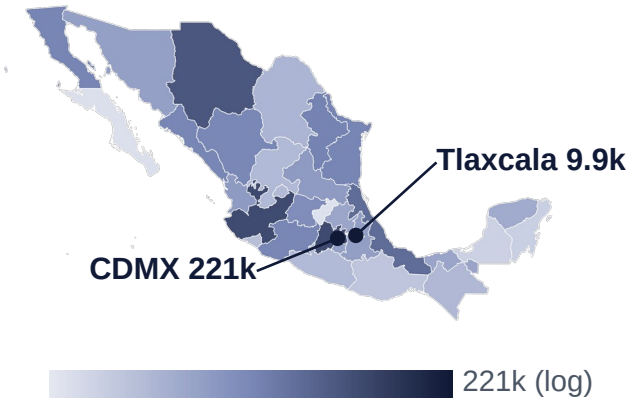
Weekly reported cases, national ICD-10 F32, SINAVE 2014–2025



~2,900
every week
2023–2025 mean

74%
are women
annual totals by sex

~22× **more cases: CDMX vs. Tlaxcala**



Source: SINAVE bulletins [1]; reported counts, not prevalence. Lines: weekly means 2017–19 and 2023–25. Map: cumulative cases 2014–2025.

02 THE FRAMEWORK

Four candidates per series, one auditable rule

SINAVE bulletins

Weekly ICD-10 F32 counts, 2014–2025 · 626 weeks

111 = (32 states + 4 regions + national) × 3 strata

99 evaluated = (32 states + national) × 3 strata

Four candidates

- **DeepAR** [3] probabilistic recurrent net
- **Prophet** [2] trend + seasonality
- **Ensemble** [4] Prophet + XGBoost blend
- **Stacking** [5, 6] Prophet, ETS, LightGBM

Deterministic selection rule

1. **Lowest historical sMAPE**
2. **Within 5% relative: lowest MASE**
3. **MASE tie: lowest RMSE**

Recorded forecast

52 weeks ahead

Audit log

which model, on what evidence

Selection evidence: per-model historical scores (paper Table 1); protocols differ by model. MASE: mean absolute scaled error [7].

Sociedad Mexicana de Inteligencia Artificial · Tecnológico de Monterrey Campus Chihuahua 50 Años · Ambassade de France au Mexique · México Francia 2026 200 Años de Amistad · Chihuahua Buró de Convenciones · ¡ah Chihuahua! · MICAI 25 Years 2000-2026

Recorded initial assignments

99 evaluated targets recorded assignment, by stratum

DeepAR Ensemble Prophet

General	Women	Men
		
32 states · DeepAR	32 states · DeepAR	32 states · DeepAR
National · Ensemble	National · Ensemble	National · Prophet

96 state targets on DeepAR · **3** national: Ensemble 2, Prophet 1

State targets assigned to DeepAR

96 / 96

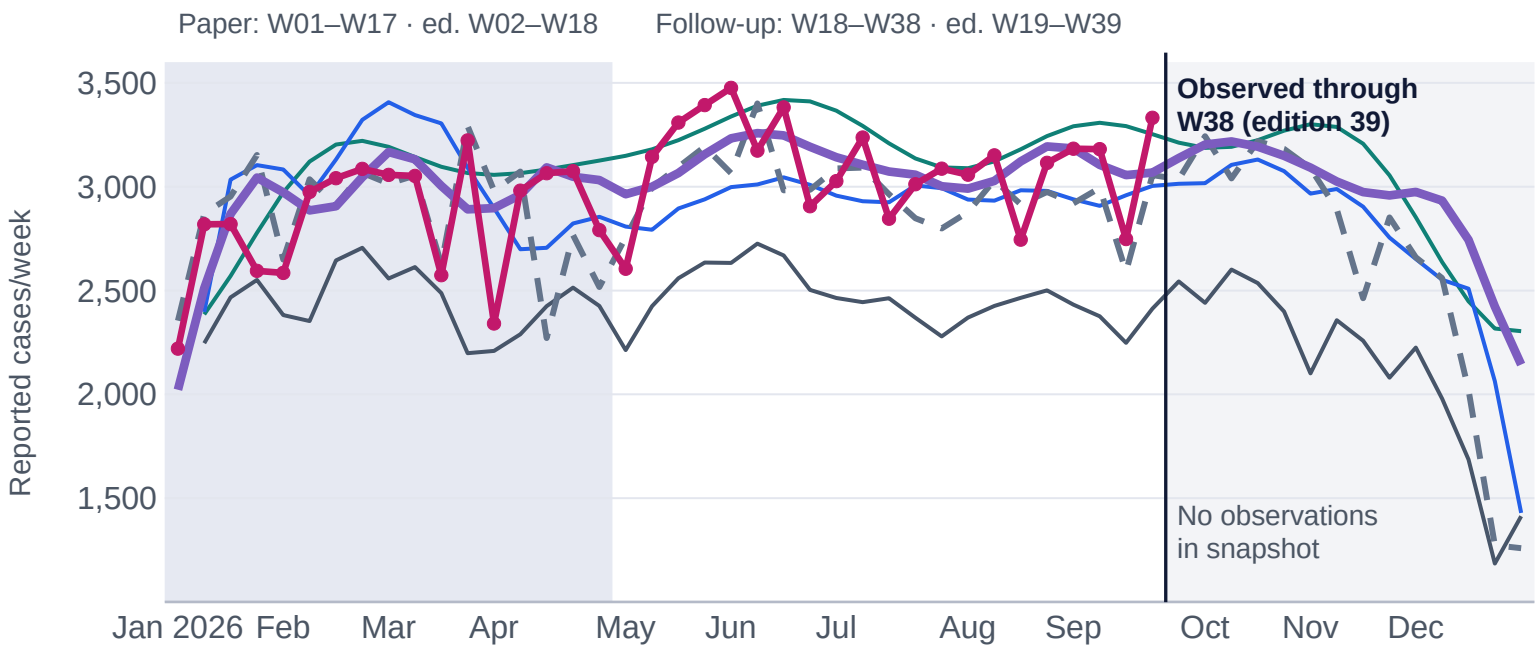
chosen from historical selection scores

Note: Historical evaluation protocols differ. The state DeepAR code scores training fit.

Assignments come from the historical selection scores of paper Table 1 (backup B8).

Recorded 2026 trajectories and follow-up

National, general stratum Prophet, DeepAR and Stacking start at W02



99 targets, editions W19–W39
median sMAPE · targets below naive

Prophet	**24.2**	72/99
Ensemble	**25.7**	70/99
DeepAR	**31.7**	57/99
Stacking	**31.4**	48/99
Archived	**31.7**	60/99
Naive	**28.6**	

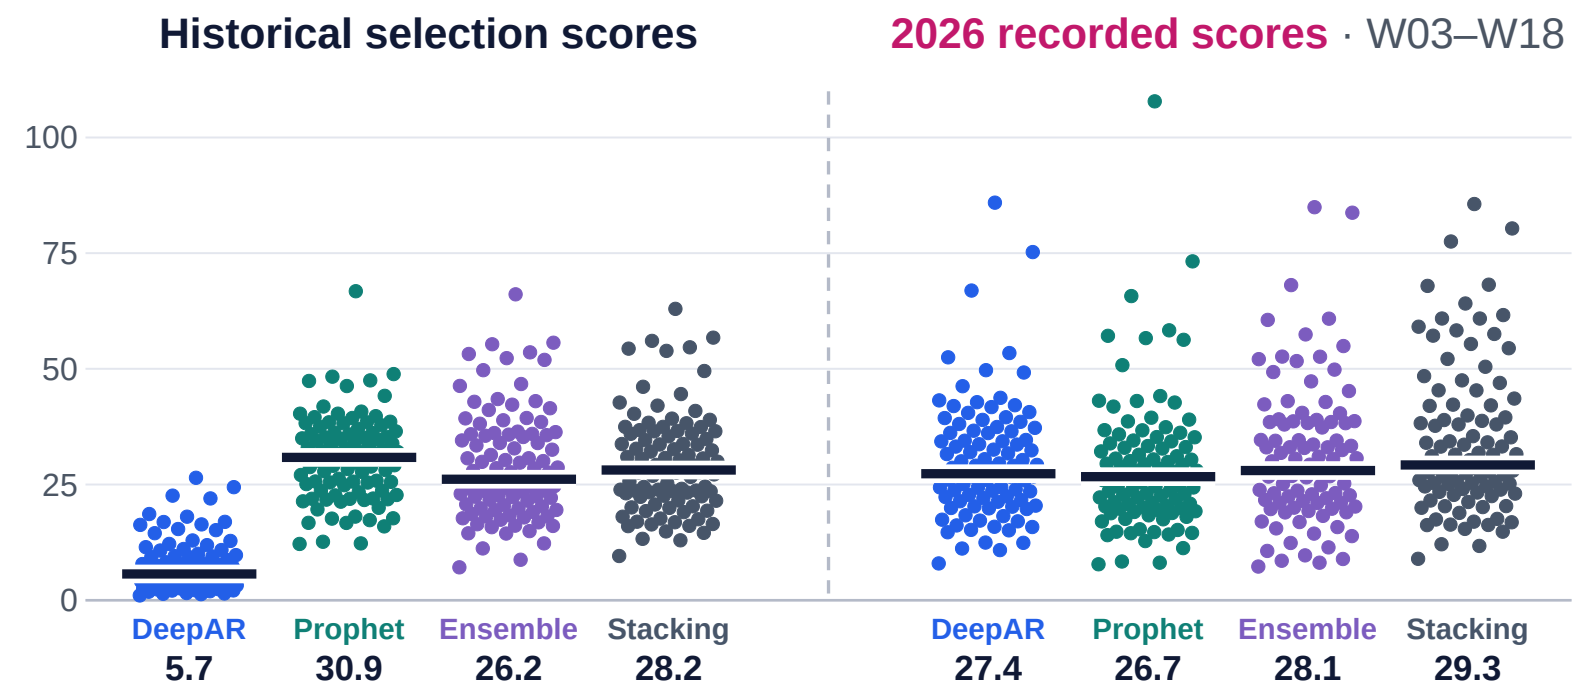
National, editions W19–W39
Lower: Ensemble

Ensemble	**5.64%**
52-week naive	**6.36%**

Recorded forecasts · retrospective evaluation. Curves: one national target; table: 99 targets. Snapshot: edition 39. Descriptive. Archived = recorded assignment.

Historical and 2026 model rankings differ

Per-series sMAPE (%) same 99 targets, one dot each · bar = median · lower is better



Historical protocols differ.
Descriptive comparison across the same 99 targets.

Range of 2026 medians

< 3 pts

Lowest median: Prophet

No model is best everywhere

35	27	19	18

targets won in 2026: Prophet, DeepAR, Ensemble, Stacking

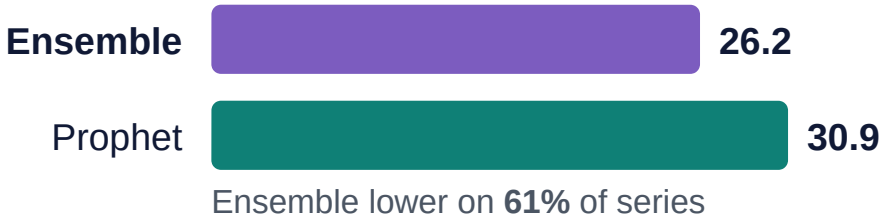
Right panel and cards: recorded errors, bulletin editions W03–W18 (paper window). Post hoc and exploratory: targets overlap; tests in B2. Full figure in B9.

06 WITHOUT DEEPAR

Without DeepAR, Prophet overtook Ensemble in W03–W18

Median sMAPE (%) Ensemble vs. Prophet · same 99 series · lower is better

Historical held-out windows



Ensemble lower on **61%** of series

2026 recorded scores W03–W18



Prophet lower on **60%** of series

Same lowest-error model, historical vs. 2026

Prophet, Ensemble, Stacking · 99 series

34% **observed**
34 of 99 series

31% **reference**
from each model's win frequencies

Descriptive reference, not a formal chance test.

Additional analysis beyond the paper. Point rankings; post hoc and exploratory. Averaging each state's three targets weakens the reversal (backup B2).

Sociedad Mexicana de Inteligencia Artificial · Tecnológico de Monterrey · Campus Chihuahua 50 Años · Ambassade de France au Mexique · México Francia 2026 200 Años de Amistad · Chihuahua Buró de Convenciones · ¡ah Chihuahua! · MICAI 25 Years 2000-2026

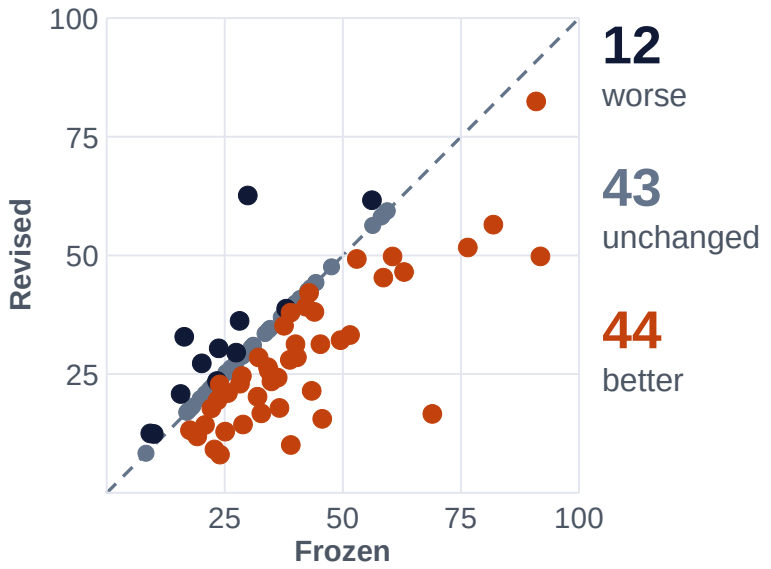
07 REVISION REPLAY

Revision lowered the median in a seven-week replay

W03–W11 re-select: same rule, prior weeks only

W12–W18 evaluation: excluded from reselection

Each dot is one series sMAPE, W12–W18



12 worse

43 unchanged

44 better

Policies median sMAPE, W12–W18

Frozen assignment	32.7
Global re-selection	32.7
Per-series revision	**28.5**
Comparators	
Prophet only	27.6
Pool without DeepAR	26.4

−4.2 pts difference of policy medians

Reading

In this seven-week replay, revision reduced the median from 32.7% to 28.5%.

Simple alternatives had lower medians.

Revised vs. frozen: 44 better, 12 worse, 43 equal. Overall medians were lower at all nine splits (W06–W14); at W11 they were lower in each stratum. Exploratory.

TAKEAWAYS

Selection is a signal, not a guarantee

An auditable per-series selection framework, evaluated on recorded 2026 forecasts.

1

Historical and 2026 rankings differ

Historical selection favored DeepAR. Prophet had the lowest median error in the recorded panel.

2

Revision lowered the replay median

Seven-week replay: revision 32.7 → 28.5; pool without DeepAR 26.4.

3

The year so far

21-week follow-up: national Ensemble 5.64 vs naive 6.36; archived panel median 31.7 vs 28.6.

THANK YOU

Questions?

Selection is a signal, not a guarantee.

Behind every weekly count are people who reached a health service.

Forecasting Weekly Depression Incidence in Mexico: A Multi-Model Framework with Auditable Per-Series Model Selection

To appear in Advances in Soft Computing (MICAI 2026), Springer LNCS 17149

rebull@exatec.tec.mx

github.com/EpiForecastMX/per-series-model-selection2026

epiforecast.mx/micai2026
follow-up, data and slides

Sociedad Mexicana de Inteligencia Artificial · Tecnológico de Monterrey · Campus Chihuahua 50 años · Ambassade de France au Mexique · México Francia 2026 · Chihuahua Buró de Convenciones · ¡ah Chihuahua! · MICAI 25 Years 2000-2026

SOURCES

References

[1] Dirección General de Epidemiología, Secretaría de Salud: SINAVE Boletín Epidemiológico (2026).

[2] Taylor, S.J., Letham, B.: Forecasting at scale. Am. Stat. 72(1), 37–45 (2018).

[3] Salinas, D., Flunkert, V., Gasthaus, J., Januschowski, T.: DeepAR: probabilistic forecasting with autoregressive recurrent networks. Int. J. Forecast. 36(3), 1181–1191 (2020).

[4] Chen, T., Guestrin, C.: XGBoost: a scalable tree boosting system. In: KDD, pp. 785–794 (2016).

[5] Ke, G., et al.: LightGBM: a highly efficient gradient boosting decision tree. In: NeurIPS 30, pp. 3146–3154 (2017).

[6] Wolpert, D.H.: Stacked generalization. Neural Netw. 5(2), 241–259 (1992).

[7] Hyndman, R.J., Koehler, A.B.: Another look at measures of forecast accuracy. Int. J. Forecast. 22(4), 679–688 (2006).

[8] Nie, Y., Nguyen, N.H., Sinthong, P., Kalagnanam, J.: A time series is worth 64 words: long-term forecasting with transformers. In: ICLR (2023).

[9] Box, G.E.P., Jenkins, G.M., Reinsel, G.C., Ljung, G.M.: Time Series Analysis: Forecasting and Control, 5th edn. Wiley (2015).

This work. Rebull-Saucedo, J.A., Pérez-Nava, J.C., Sánchez-Salazar, L.G., Barceló-Alonso, G., Pérez-Hernández, R.: Forecasting weekly depression incidence in Mexico: a multi-model framework with auditable per-series model selection. In: Advances in Soft Computing, MICAI 2026, LNCS 17149, Springer (to appear).

The selection rule, exactly

Per series, with recorded historical sMAPE, MASE, and RMSE per model:

1. If the second-lowest sMAPE exceeds 1.05 times the lowest, choose the lowest.
2. Otherwise, among models with sMAPE ≤ 1.05 × best, pick the lowest MASE.
3. If MASE is tied, pick the lowest RMSE.

\+ Fewer than 5 cases in the 52 weeks before the forecast origin: use the regional model (past data only; never triggered here, 0 of 111).

Each choice and its metrics are stored with the forecast: the audit log.

Selection evidence

Per-model historical backtests (paper Table 1), not one shared protocol.

Prophet, Ensemble, Stacking: held-out windows, with different protocols.

Historical state DeepAR code skips CV, trains on the full series, then scores that fit (commit 60d6ea0f; changed in 9fdd6dca). The exact published runs are not independently linked.

The paper reports a data cutoff of 2025-W53; execution dates are not independently verified.

Statistical tests and uncertainty

Comparison	99 targets: raw / Holm	32 states: nominal
Historical, Ensemble vs. Prophet	0.020 / 0.060	0.15
2026, Prophet vs. Ensemble	0.007 / 0.035	0.06
2026, Prophet vs. DeepAR	0.070 / 0.14	0.24
Revision vs. frozen, split W11	< 0.001 / < 0.001	< 0.001

Exploratory, post hoc; targets overlap and Holm does not remove dependence. Not significant ≠ equivalent. Two-sided Wilcoxon, zsplit; Holm within six model pairs (nine for revision).

National, W03–W18, recomputed

Mean Poisson deviance: Ensemble 25.7, Prophet 32.3, DeepAR 46.0, Stacking 97.8.

Reported in the paper only

Diebold–Mariano tests; the 80% band covered 10 of 17 weeks. Inputs not released.

Limitations

One condition

Depression (F32) only; transfer is untested.

Training-fit evidence

Historical state DeepAR code uses training fit; run provenance is unverified.

Unequal search budget

Prophet had the broadest search; tuning effect untested.

Partial-year evaluation

Paper windows plus 21 follow-up periods (to W38); full year still pending.

Provenance

Forecast execution dates not independently verified.

Uncertainty

Paper-reported 80% band coverage: 58.8%; revision does not calibrate it.

Reported incidence

Models learn reporting and disease together.

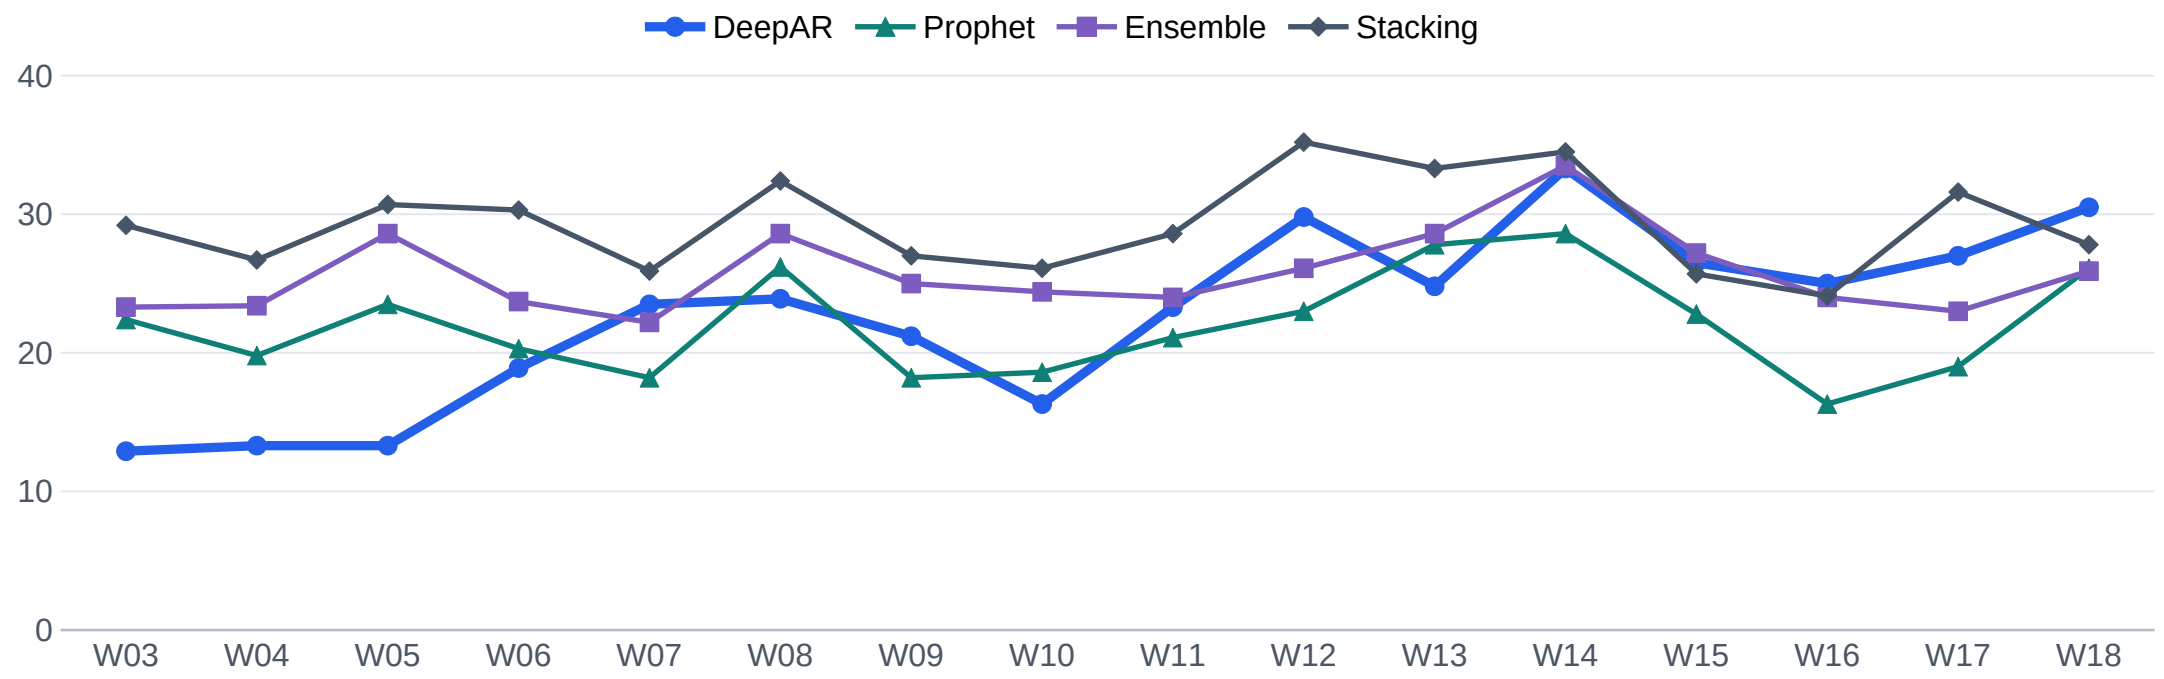
No external validation

Only SINAVE / Mexico; a research prototype.

B4 BACKUP · HORIZON

DeepAR leads the first four editions (W03–W06)

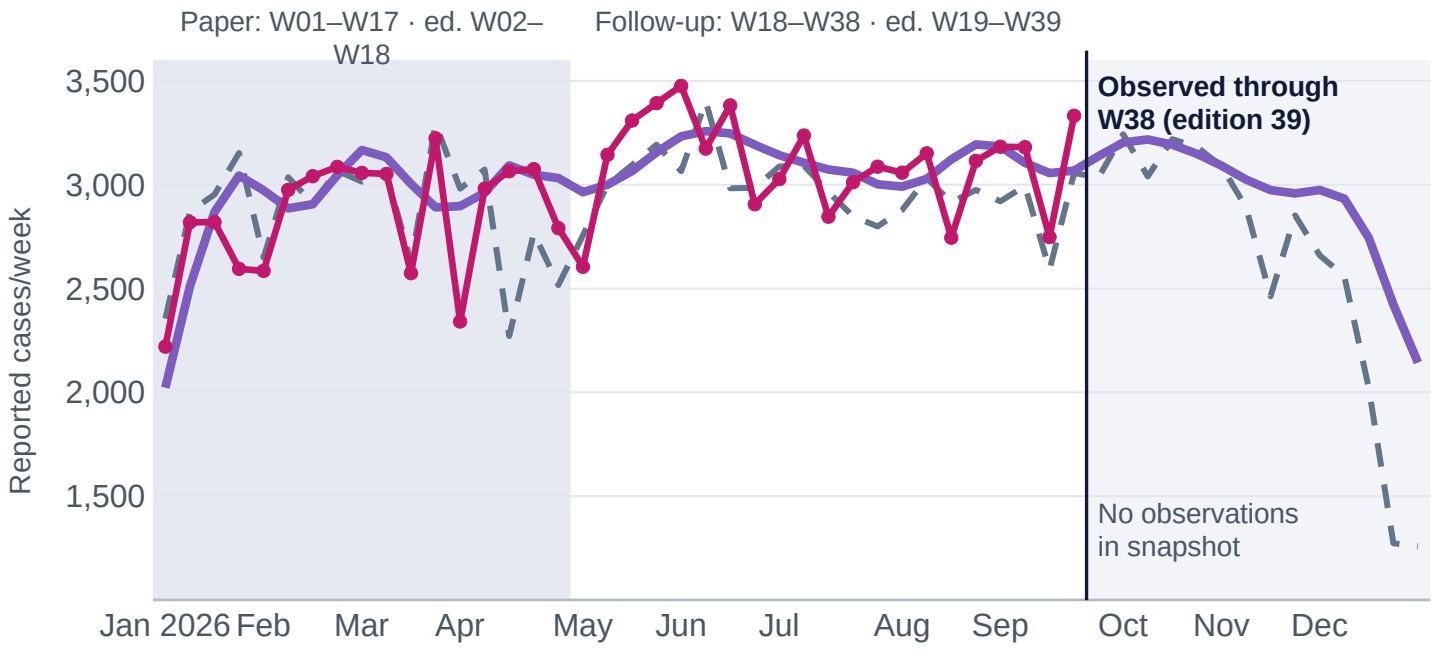
Median per-week sMAPE (%) across the 99 series 2026 bulletin editions · lower is better



Revision chooses among frozen 52-week paths; horizon effects may contribute to its gain (not isolated). Recomputed from the public files.

B5 BACKUP · NATIONAL

National forecast versus the 52-week naive



Paper sMAPE (%)

editions W02–W18 · **Lower: naive**

Ensemble	**7.40%**
52-week naive	**7.33%**

Follow-up sMAPE (%)

editions W19–W39 · **Lower: Ensemble**

Ensemble	**5.64%**
52-week naive	**6.36%**

Forecast periods mapped by rule; execution origin unverified.

sMAPE, lower is better. Cumulative deviation, Ensemble / naive: paper +2.45% / +0.6%; follow-up +0.21% / −3.81%. Snapshot: edition 39. Descriptive.

Sociedad Mexicana de Inteligencia Artificial · Tecnológico de Monterrey · Campus Chihuahua 50 Años · Ambassade de France au Mexique · México Francia 2026 · Chihuahua Buró de Convenciones · ¡ah Chihuahua! · MICAI 25 Years 2000-2026

B6 BACKUP · BY SEX

Every model against the naive, by window and sex

Median sMAPE (%) and targets beating the naive, of 33 per stratum · lower is better

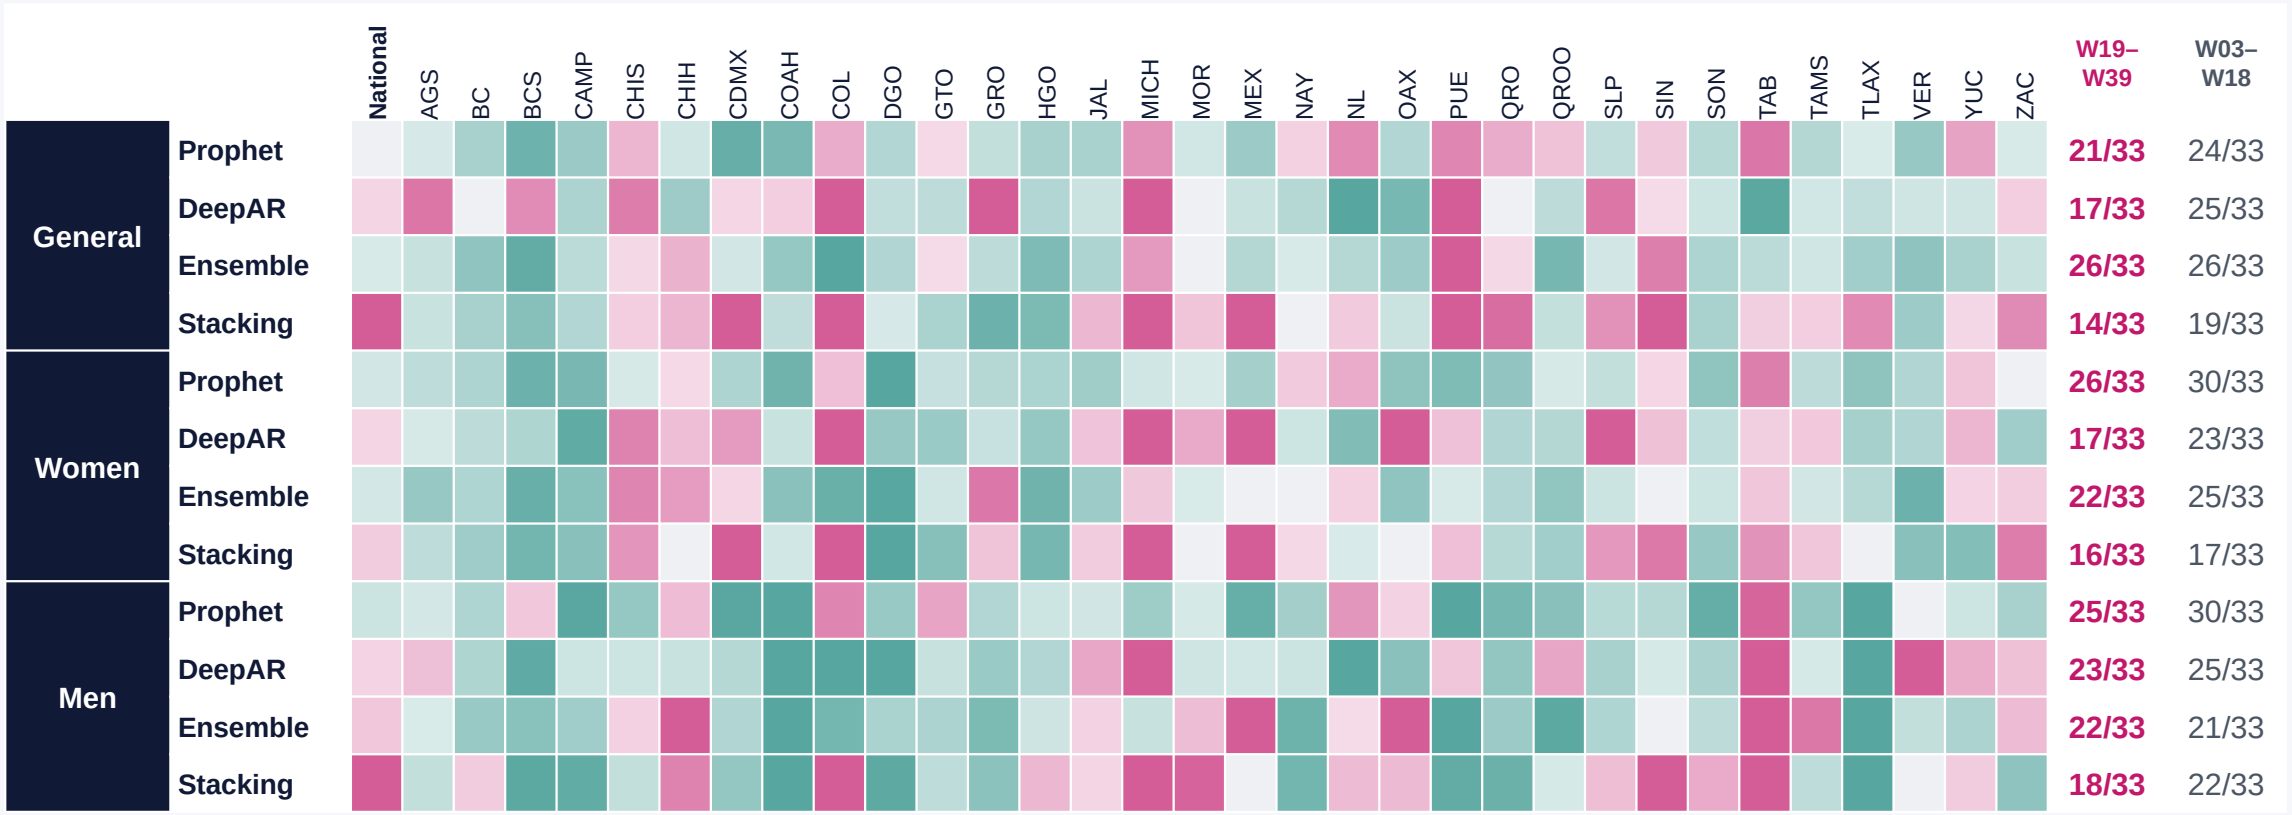
	Paper window · editions W03–W18			Follow-up · editions W19–W39		
Estimator	**General**	**Women**	**Men**	**General**	**Women**	**Men**
Prophet	**23.8 · 24**	24.4 · 30	**30.4 · 30**	**21.3 · 21**	**22.7 · 26**	**30.0 · 25**
DeepAR	25.9 · 25	26.0 · 23	32.4 · 25	23.3 · 17	27.1 · 17	35.2 · 23
Ensemble	26.7 · 26	**23.8 · 25**	38.0 · 21	21.8 · 26	25.7 · 22	36.3 · 22
Stacking	29.0 · 19	27.9 · 17	34.3 · 22	27.8+ · 14	30.6+ · 16	39.7+ · 18
Archived assignment	25.9 · 26	26.0 · 24	32.4 · 25	23.3 · 18	27.1 · 18	35.2 · 24
52-week naive	31.2	30.6	40.1	23.7	27.8	39.0

Cell: median sMAPE across 33 targets · targets with lower sMAPE than the naive. + = median above the naive; bold = lowest median. Archived assignment = DeepAR in every state. Descriptive; overlapping targets.

B7 BACKUP · BY STATE

How the year is going: every model against the naive

Follow-up, editions W19–W39 · cell = model sMAPE minus naive sMAPE, one target each



Model minus naive (pts): ≤−15 · −8 · −3 · ±0.5 · +3 · +8 · ≥+15

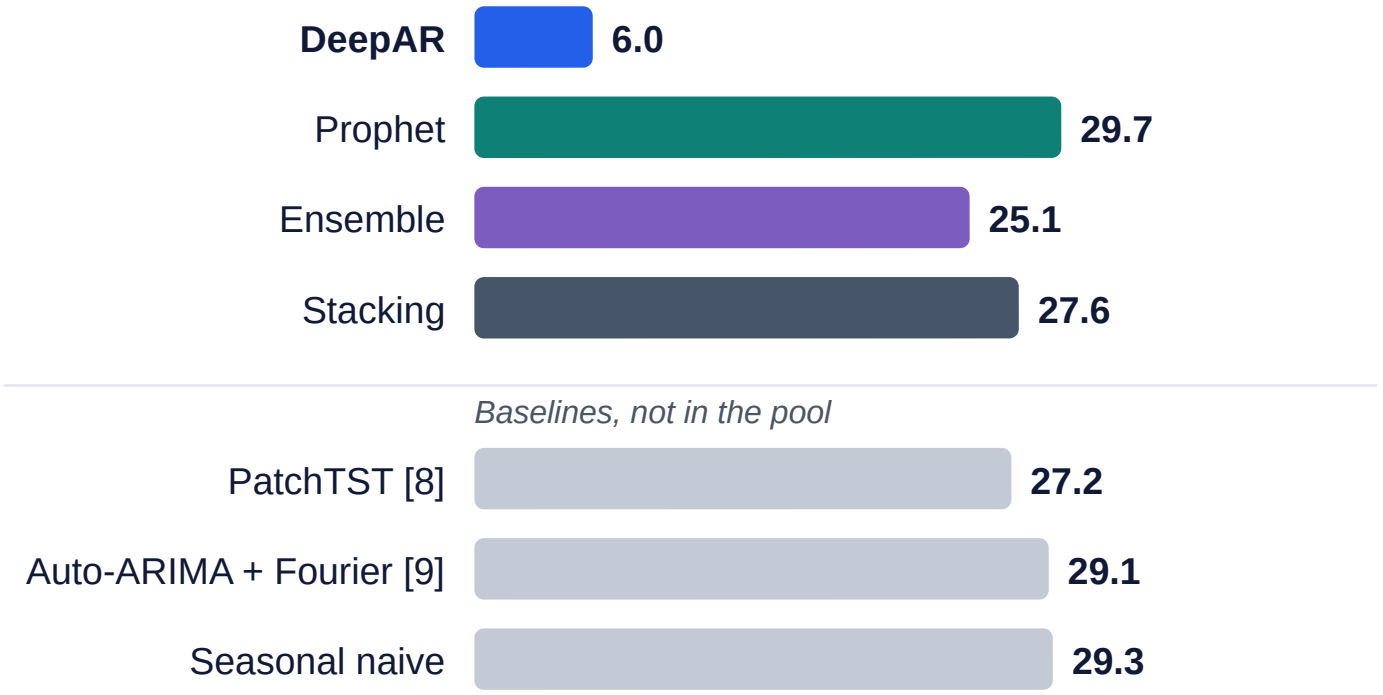
Teal: lower error than the naive (better). Magenta: higher. Gray: under 0.5 pts, not necessarily a tie. Archived assignment = DeepAR in every state. Descriptive; overlapping targets. Full matrix for download; state-level reading requires zoom.

Sociedad Mexicana de Inteligencia Artificial · Tecnológico de Monterrey · Campus Chihuahua 50 Años · Ambassade de France au Mexique · México Francia 2026 · Chihuahua Buró de Convenciones · ¡ah Chihuahua! · MICAI 25 Years 2000-2026

B8 BACKUP · HISTORICAL SELECTION

Historical selection scores (paper Table 1)

Median historical sMAPE (%) Paper Table 1 · 111 targets · lower is better



Series the rule assigns to DeepAR

107 / 111

96.4%, including every state series

Reading the 6.0: in the historical code, DeepAR's state scores reflect training fit, not forecast error.

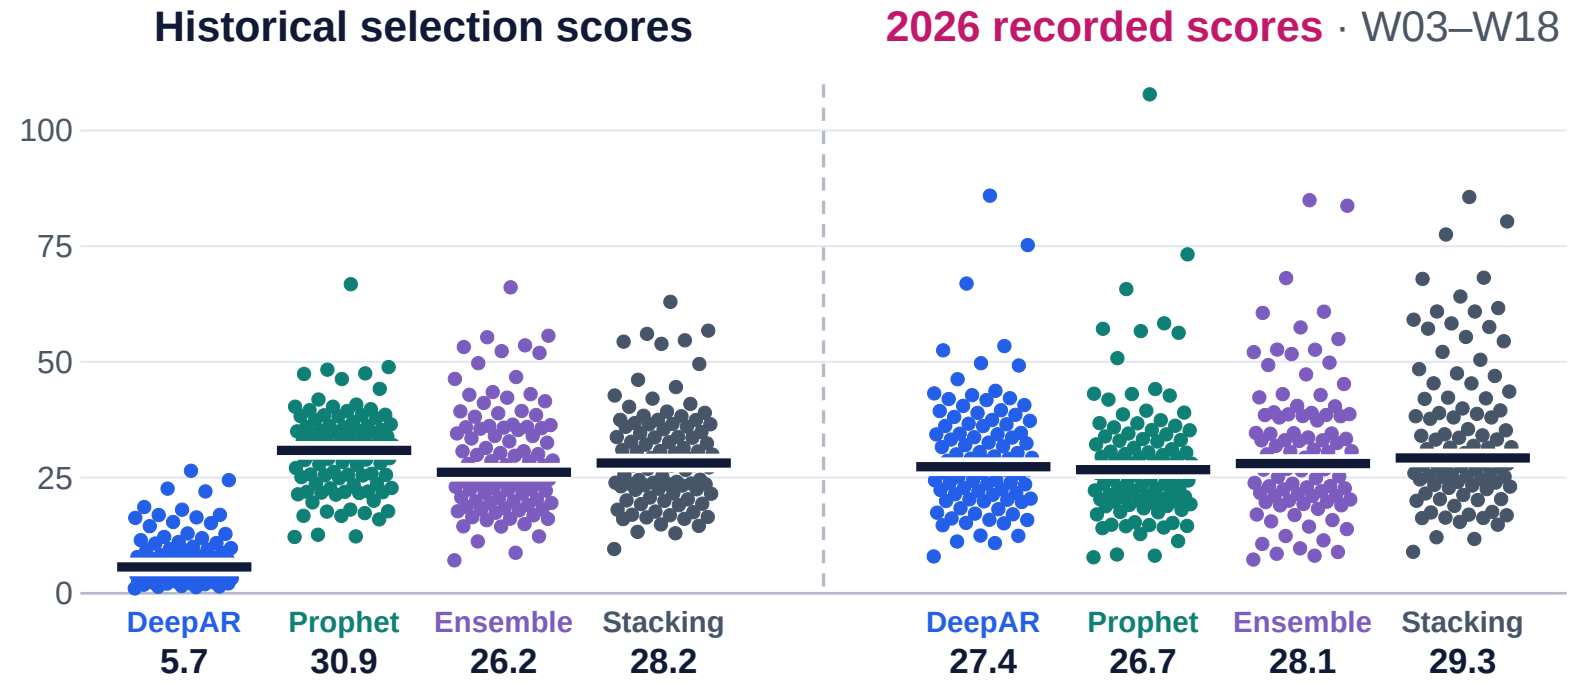
DeepAR's 6.0% historical sMAPE is the median over all 111 targets. Historical selection scores are not a guarantee of 2026 accuracy.

Sociedad Mexicana de Inteligencia Artificial · Tecnológico de Monterrey · Campus Chihuahua 50 Años · Ambassade de France au Mexique · México Francia 2026 · Chihuahua Buró de Convenciones · ¡ah Chihuahua! · MICAI 25 Years 2000-2026

B9 BACKUP · HISTORICAL VS. 2026

Historical and 2026 model rankings differ

Per-series sMAPE (%) 99 series, one dot each · bar = median · lower is better



Historical score gap

≈20 pts

Range of 2026 medians

< 3 pts

No model is best everywhere

35 | 27 | 19 | 18

series won in 2026: Prophet, DeepAR, Ensemble, Stacking

Historical protocols differ by model. Post hoc and exploratory: targets overlap; paired tests in backup B2.

Sociedad Mexicana de Inteligencia Artificial · Tecnológico de Monterrey Campus Chihuahua 50 Años · Ambassade de France au Mexique · México Francia 2026 · Chihuahua Buró de Convenciones · ¡ah Chihuahua! · MICAI 25 Years 2000-2026