MICAI 25 YEARS 2000-2026
02-06 NOV 2026
MEXICAN INTERNATIONAL CONFERENCE ON ARTIFICIAL INTELLIGENCE
|ITESM Campus Chihuahua

Paper #12 · Oral

Forecasting Weekly Depression Incidence in Mexico:
A Multi-Model Framework with Auditable Per-Series Model Selection

Javier Augusto Rebull-Saucedo[1] · **Juan Carlos Pérez-Nava**[1] · **Luis Gerardo Sánchez-Salazar**[1]
Grettel Barceló-Alonso[1] · **Ruth Pérez-Hernández**[2]

[1] Tecnológico de Monterrey, Monterrey, Mexico
[2] Instituto Mexicano del Seguro Social, OOAD Guerrero, Mexico
Independent research; affiliations do not imply institutional endorsement.

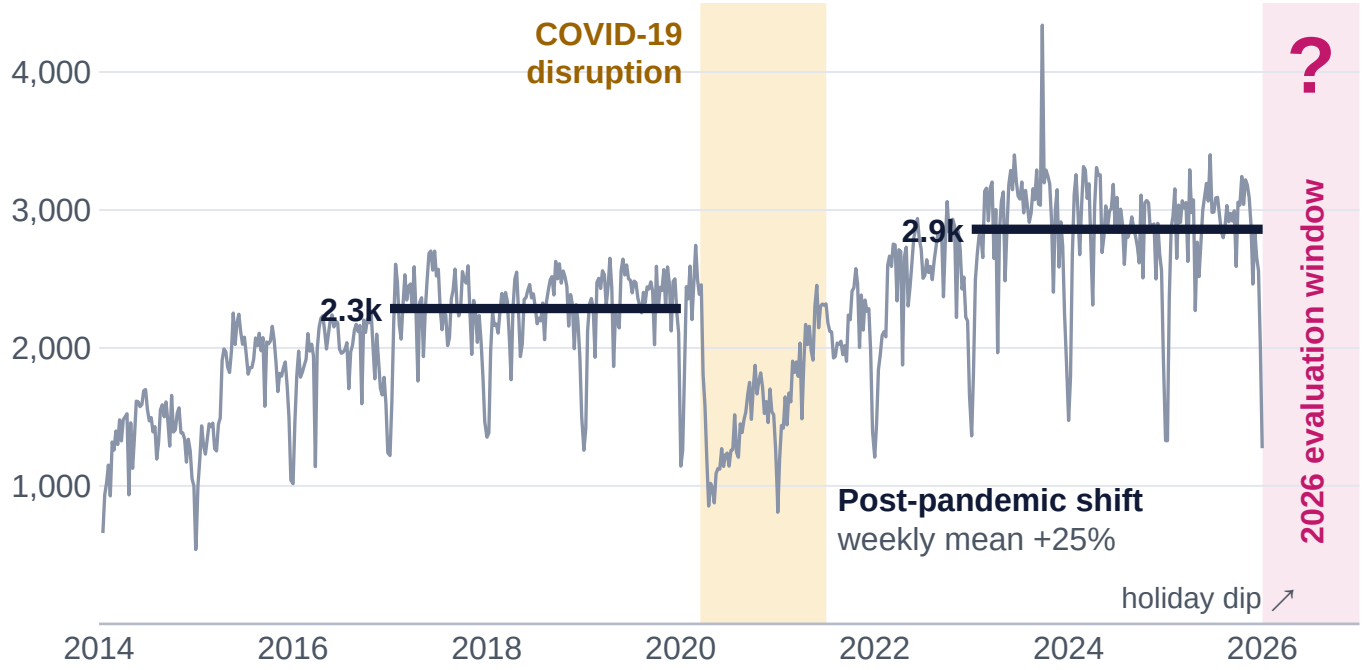
When historical scores meet recorded 2026 forecasts

Sociedad Mexicana de Inteligencia Artificial · Tecnológico de Monterrey Campus Chihuahua 50 años · Ambassade de France au Mexique · México Francia 2026 · Chihuahua Buró de Convenciones · ¡ah Chihuahua!

01 THE PROBLEM

Depression is counted weekly, but only in hindsight

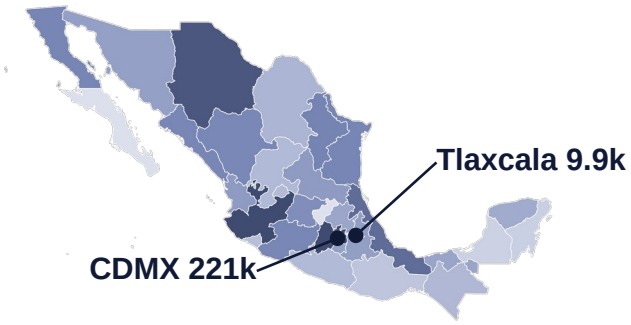
Weekly reported cases, national ICD-10 F32, SINAVE 2014–2025



~150k
cases per year
2023–2025

74%
are women
annual totals by sex

~22× **spread across states**



9.9k — 221k (log)

Source: SINAVE bulletins [1]; reported counts, not prevalence. Lines: weekly means 2017–19, 2023–25. Curve omits 2014-W02 and 2016-W20 (flagged partial).

Sociedad Mexicana de Inteligencia Artificial · Tecnológico de Monterrey Campus Chihuahua 50 Años · Ambassade de France au Mexique · México Francia 2026 · Chihuahua Buró de Convenciones · ¡ah Chihuahua! · MICAI 25 Years 2000-2026

02 THE FRAMEWORK

Four candidates per series, one auditable rule

SINAVE bulletins

Weekly ICD-10 F32 counts, 2014–2025 (626 weeks)

111 targets: 32 states, 4 regions, national; general, women, men

99 reported directly

Four candidates

- **DeepAR** [3] probabilistic recurrent net
- **Prophet** [2] trend + seasonality
- **Ensemble** [4] Prophet + XGBoost blend
- **Stacking** [5, 6] Prophet, ETS, LightGBM

Deterministic selection rule

1. **Lowest historical sMAPE**
2. **Within 5% relative: lowest MASE**
3. **MASE tie: lowest RMSE**

Recorded forecast

52 weeks ahead

Audit log

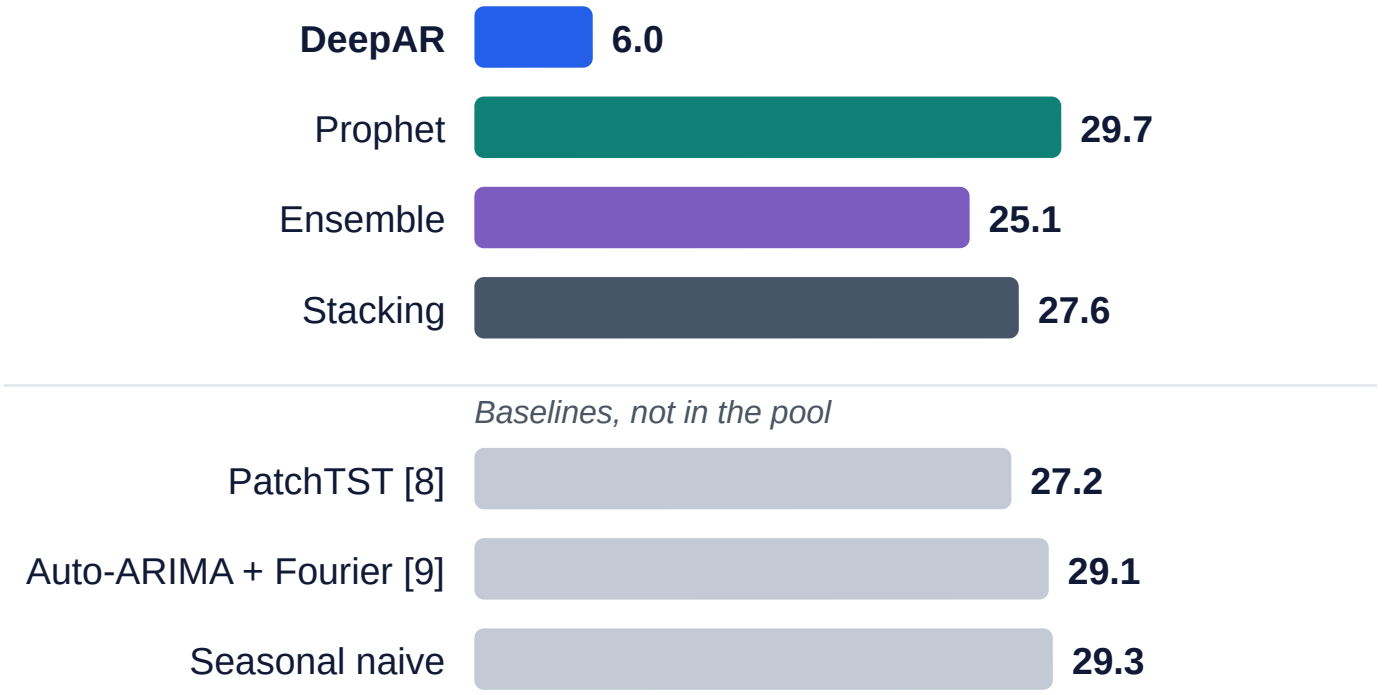
which model, on what evidence

Selection evidence: per-model historical scores (paper Table 1); protocols differ by model. MASE as in [7]. The sparse fallback never triggered (0 of 111).

03 SELECTION EVIDENCE

Historical selection favored DeepAR

Median historical sMAPE (%) Paper Table 1 · 111 targets · lower is better



Series the rule assigns to DeepAR

107 / 111

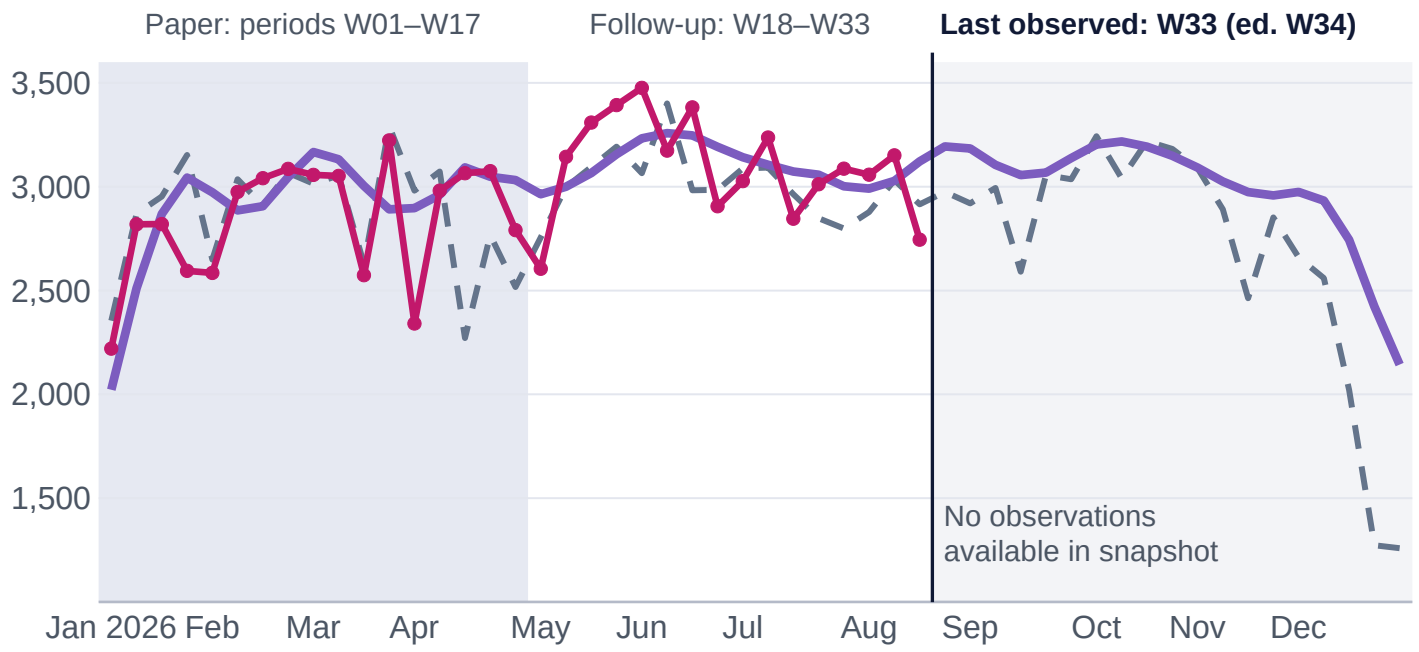
96.4%, including every state series

Caveat: Historical state DeepAR code scores training fit; a common validation protocol is not verified.

DeepAR's 6.0% historical sMAPE is the median over all 111 targets. Historical selection scores are not a guarantee of 2026 accuracy.

Sociedad Mexicana de Inteligencia Artificial · Tecnológico de Monterrey Campus Chihuahua 50 Años · Ambassade de France au Mexique · México Francia 2026 200 Años de Amistad · Chihuahua Buró de Convenciones · ¡ah Chihuahua! · MICAI 25 Years 2000-2026

Recorded trajectories and available follow-up



Paper editions W02–W18	sMAPE (%)
Ensemble	7.40%
52-week naive	7.33%

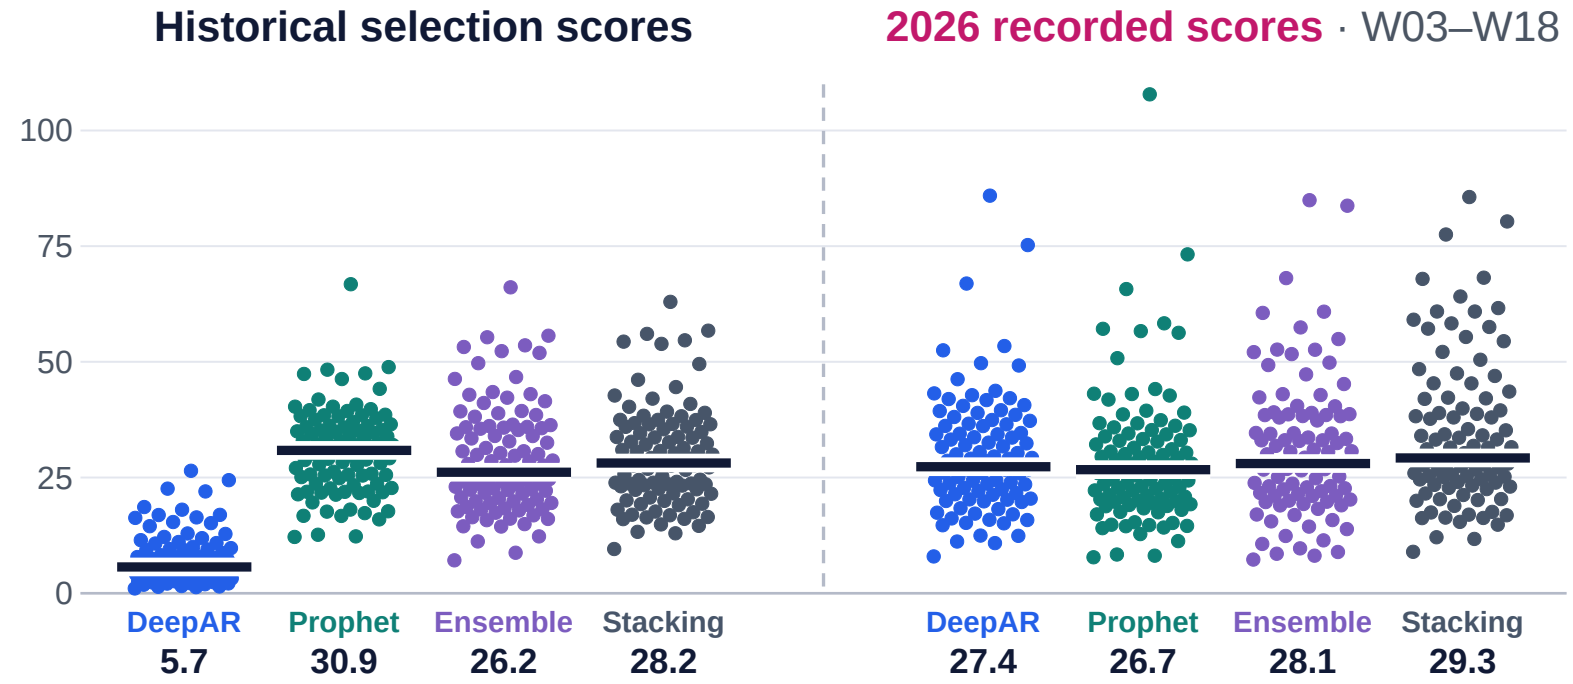
Follow-up editions W19–W34	sMAPE (%)
Ensemble	5.91%
52-week naive	6.24%

Recorded forecasts · retrospective evaluation

sMAPE. Cumulative deviation, Ensemble / naive: paper +2.45% / +0.6%; follow-up +0.17% / −2.95%. Snapshot: edition W34. Descriptive.

Historical and 2026 model rankings differ

Per-series sMAPE (%) 99 series, one dot each · bar = median · lower is better



DeepAR's historical margin

≈20 pts

2026 spread, all four

< 3 pts

No model is best everywhere



series won in 2026: Prophet, DeepAR, Ensemble, Stacking

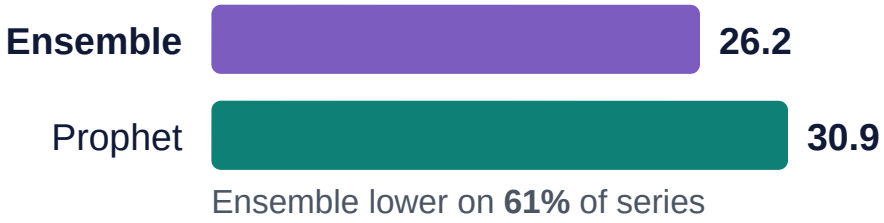
Historical protocols differ by model. Post hoc and exploratory: targets overlap; paired tests in backup B2.

06 WITHOUT DEEPAR

Excluding DeepAR, the two rankings differ

Median sMAPE (%) Ensemble vs. Prophet · same 99 series · lower is better

Historical held-out windows



Ensemble lower on **61%** of series

2026 recorded scores W03–W18



Prophet lower on **60%** of series

Same lowest-error model, historical vs. 2026

Prophet, Ensemble, Stacking · 99 series

34% **observed**
34 of 99 series

31% **reference**
from each model's win frequencies

Descriptive reference, not a formal chance test.

Point rankings; post hoc and exploratory. Averaging each state's three targets weakens the reversal (backup B2).

Sociedad Mexicana de Inteligencia Artificial · Tecnológico de Monterrey · Campus Chihuahua 50 Años · Ambassade de France au Mexique · México Francia 2026 · Chihuahua Buró de Convenciones · ¡ah Chihuahua! · MICAI 25 Years 2000-2026

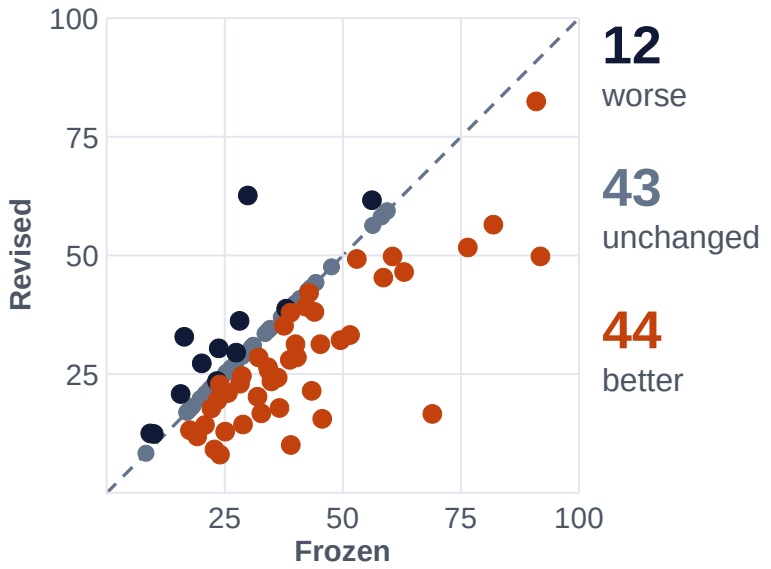
07 REVISION REPLAY

Per-series revision beat the frozen and global picks

W03–W11 re-select: same rule, prior weeks only

W12–W18 evaluation: excluded from reselection

Each dot is one series sMAPE, W12–W18



12 worse

43 unchanged

44 better

Policies median sMAPE, W12–W18

Frozen assignment	32.7
Global re-selection	32.7
Per-series revision	**28.5**
Comparators	
Prophet only	27.6
Pool without DeepAR	26.4

−4.2 pts difference of policy medians

Reading

In this seven-week replay, revision reduced the median from 32.7% to 28.5%.

Simple alternatives had lower medians.

Revised vs. frozen: 44 better, 12 worse, 43 equal; lower medians at every split W06–W14 and in each stratum (exploratory).

Sociedad Mexicana de Inteligencia Artificial · Tecnológico de Monterrey · Campus Chihuahua 50 años · Ambassade de France au Mexique · México Francia 2026 · Chihuahua Buró de Convenciones · ¡ah Chihuahua! · MICAI 25 Years 2000-2026

TAKEAWAYS

Selection is a signal, not a guarantee

An auditable per-series selection framework, evaluated on recorded 2026 forecasts.

1

Historical and 2026 rankings differ

Historical selection did not identify the lowest-error model in the recorded 2026 panel.

2

Revision lowered the replay median

Seven-week replay: revision 32.7 → 28.5; pool without DeepAR 26.4.

3

The year so far

16-week follow-up: national Ensemble 5.91 vs naive 6.24; archived panel median 31.0 vs 27.0.

THANK YOU

Questions?

Selection is a signal, not a guarantee.

Forecasting Weekly Depression Incidence in Mexico: A Multi-Model Framework with Auditable Per-Series Model Selection

To appear in Advances in Soft Computing (MICAI 2026), Springer LNCS 17149

rebull@exatec.tec.mx

github.com/EpiForecastMX/per-series-model-selection2026

epiforecast.mx/micai2026
follow-up, data and slides

Sociedad Mexicana de Inteligencia Artificial
Tecnológico de Monterrey | CAMPUS CHIHUAHUA 50 AÑOS
AMBASSADE DE FRANCE AU MEXIQUE
MÉXICO FRANCIA 2026 200 AÑOS DE AMISTAD
CHIHUAHUA BURÓ DE CONVENCIONES
¡ah Chihuahua!
MICAI 25 YEARS 2000-2026

SOURCES

References

[1] Dirección General de Epidemiología, Secretaría de Salud: SINAVE Boletín Epidemiológico (2026).

[2] Taylor, S.J., Letham, B.: Forecasting at scale. Am. Stat. 72(1), 37–45 (2018).

[3] Salinas, D., Flunkert, V., Gasthaus, J., Januschowski, T.: DeepAR: probabilistic forecasting with autoregressive recurrent networks. Int. J. Forecast. 36(3), 1181–1191 (2020).

[4] Chen, T., Guestrin, C.: XGBoost: a scalable tree boosting system. In: KDD, pp. 785–794 (2016).

[5] Ke, G., et al.: LightGBM: a highly efficient gradient boosting decision tree. In: NeurIPS 30, pp. 3146–3154 (2017).

[6] Wolpert, D.H.: Stacked generalization. Neural Netw. 5(2), 241–259 (1992).

[7] Hyndman, R.J., Koehler, A.B.: Another look at measures of forecast accuracy. Int. J. Forecast. 22(4), 679–688 (2006).

[8] Nie, Y., Nguyen, N.H., Sinthong, P., Kalagnanam, J.: A time series is worth 64 words: long-term forecasting with transformers. In: ICLR (2023).

[9] Box, G.E.P., Jenkins, G.M., Reinsel, G.C., Ljung, G.M.: Time Series Analysis: Forecasting and Control, 5th edn. Wiley (2015).

This work. Rebull-Saucedo, J.A., Pérez-Nava, J.C., Sánchez-Salazar, L.G., Barceló-Alonso, G., Pérez-Hernández, R.: Forecasting weekly depression incidence in Mexico: a multi-model framework with auditable per-series model selection. In: Advances in Soft Computing, MICAI 2026, LNCS 17149, Springer (to appear).

The selection rule, exactly

Per series, with recorded historical sMAPE, MASE, and RMSE per model:

1. If the best sMAPE beats the second by more than 5% (relative), pick it.
2. Otherwise, among models with sMAPE ≤ 1.05 × best, pick the lowest MASE.
3. If MASE is tied, pick the lowest RMSE.

\+ Fewer than 5 cases in the 52 weeks before the origin: use the regional model (past data only; never triggered here, 0 of 111).

Each choice and its metrics are stored with the forecast: the audit log.

Selection evidence

Per-model historical backtests (paper Table 1), not one shared protocol.

Prophet, Ensemble, Stacking: held-out windows, with different protocols.

Historical state DeepAR code skips CV, trains on the full series, then scores that fit (commit 60d6ea0f; changed in 9fdd6dca). The exact published runs are not independently linked.

The paper reports a data cutoff of 2025-W53; execution dates are not independently verified.

Statistical tests and uncertainty

Comparison	99 targets: raw / Holm	32 states: nominal
Historical, Ensemble vs. Prophet	0.020 / 0.060	0.15
2026, Prophet vs. Ensemble	0.007 / 0.035	0.06
2026, Prophet vs. DeepAR	0.070 / 0.14	0.24
Revision vs. frozen, split W11	< 0.001 / < 0.001	< 0.001

Exploratory, post hoc; targets overlap and Holm does not remove dependence. Not significant ≠ equivalent. Two-sided Wilcoxon, zsplit; Holm within six model pairs (nine for revision).

National, W03–W18, recomputed

Mean Poisson deviance: Ensemble 25.7, Prophet 32.3, DeepAR 46.0, Stacking 97.8.

Reported in the paper only

Diebold–Mariano tests; the 80% band covered 10 of 17 weeks. Inputs not released.

Limitations

One condition

Depression (F32) only; transfer is untested.

Training-fit evidence

Historical state DeepAR code uses training fit; run provenance is unverified.

Unequal search budget

Prophet had the broadest search; tuning effect untested.

Partial-year evaluation

Paper windows plus 16 follow-up periods; full-year evaluation remains pending.

Provenance

Forecast execution dates not independently verified.

Uncertainty

Paper-reported 80% band coverage: 58.8%; revision does not calibrate it.

Reported incidence

Models learn reporting and disease together.

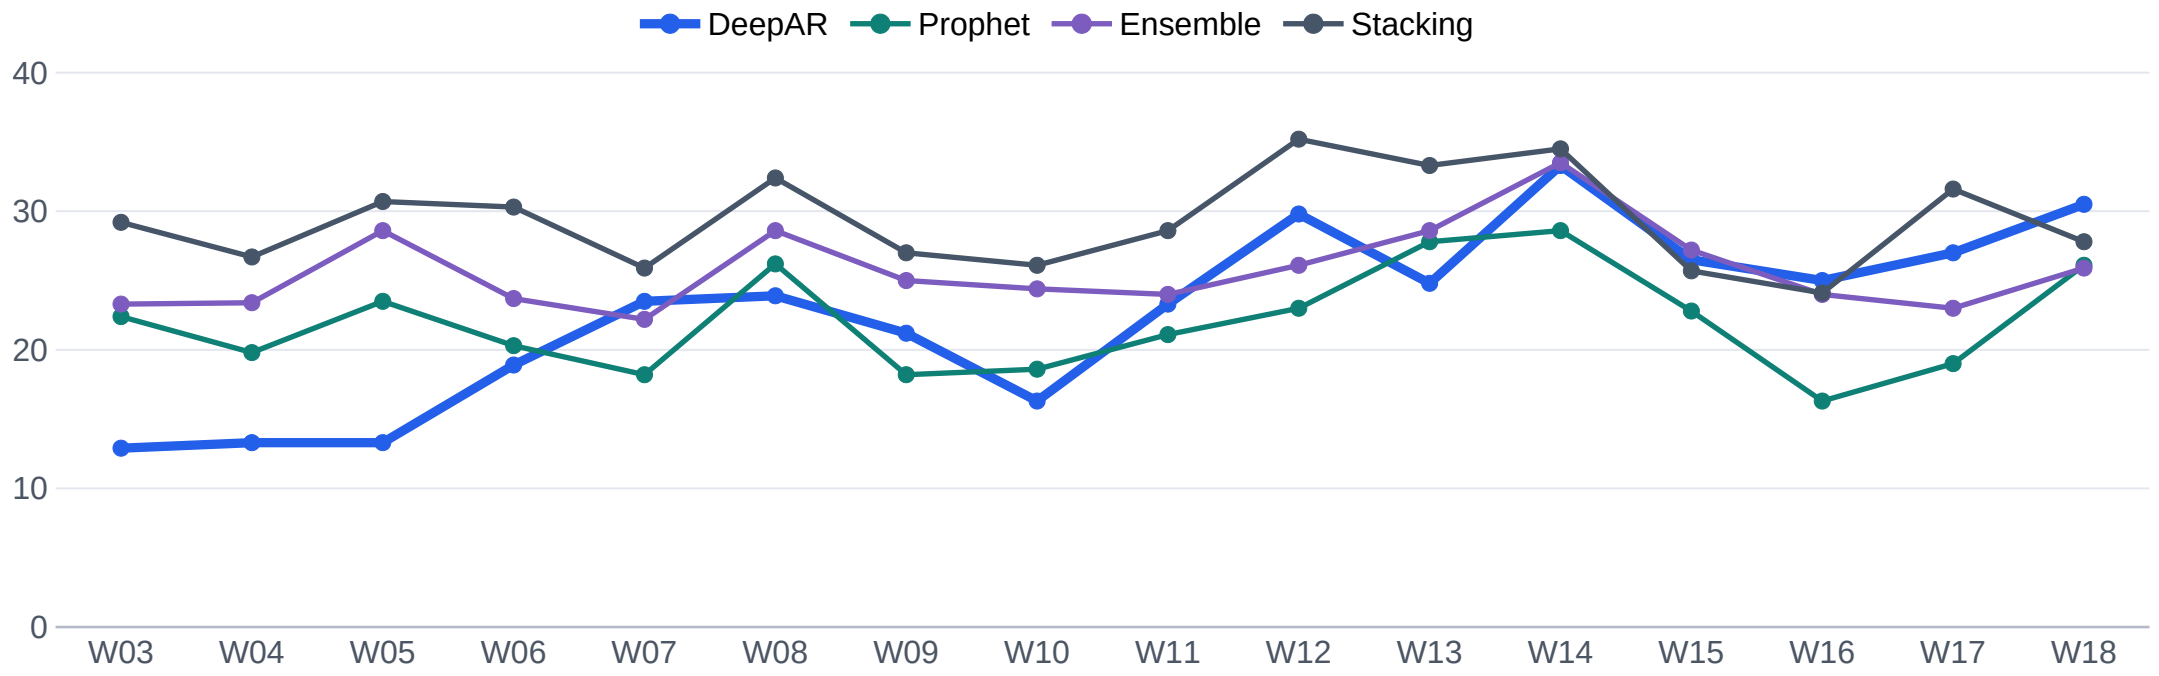
No external validation

Only SINAVE / Mexico; a research prototype.

B4 BACKUP · HORIZON

DeepAR leads the first four weeks; later ranks vary

Median per-week sMAPE (%) across the 99 series 2026 bulletin weeks · lower is better



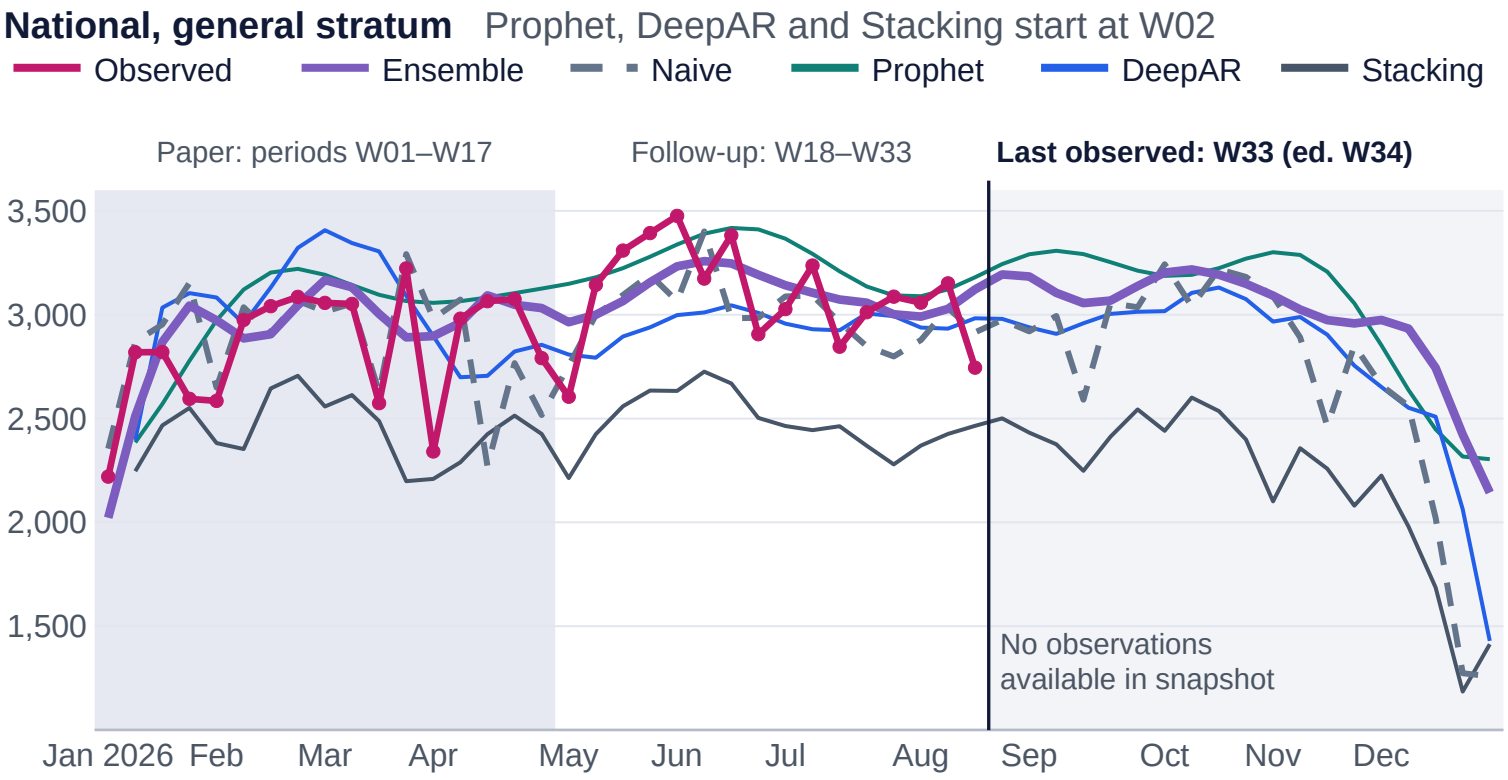
Revision chooses among frozen 52-week paths; horizon effects may contribute to its gain (not isolated). Recomputed from the public files.

Sociedad Mexicana de Inteligencia Artificial
Tecnológico de Monterrey
CAMPUS CHIHUAHUA 50 AÑOS
AMBASSADE DE FRANCE AU MEXIQUE
MÉXICO FRANCIA 2026 200 AÑOS DE AMISTAD
CHIHUAHUA BURÓ DE CONVENCIONES
¡ah Chihuahua!
MICAI 25 YEARS 2000-2026

B5 BACKUP · FOLLOW-UP

Recorded 2026 paths: all four engines

National, general stratum Prophet, DeepAR and Stacking start at W02



99 targets, editions W19–W34
median sMAPE · targets below naive

Prophet	**23.4**	70/99
Ensemble	**25.2**	71/99
DeepAR	**31.0**	52/99
Stacking	**30.6**	51/99
Archived	**31.0**	55/99
Naive	**27.0**	

Forecast periods mapped by rule; execution origin unverified.

Descriptive; targets overlap. W01 reproduces the paper; legacy mapping is ambiguous. Archived = frozen assignment (96 DeepAR).